The background of the slide features a large, faint, red watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center, with the words "STANFORD UNIVERSITY" and "1891" around it. The text "FREIHEIT WEHT" is also visible at the top of the seal.

Lecture 4: Pixels and Filters

Dr. Juan Carlos Niebles
Stanford AI Lab

Professor Fei-Fei Li
Stanford Vision Lab

What we will learn today?


- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

Images as functions

- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
(or “intensity”): $[0, 255]$



Images as functions

- An image contains discrete number of pixels

- A simple example

- Pixel value:

- “grayscale”


- (or “intensity”): [0,255]

- “color”

- RGB: [R, G, B]

- Lab: [L, a, b]

- HSV: [H, S, V]




Images as functions

- **An Image** as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$f: [a,b] \times [c,d] \rightarrow [0,255]$$

Domain
support

range



Images as functions

- **An Image** as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$f: \underbrace{[a, b] \times [c, d]}_{\substack{\text{Domain} \\ \text{support}}} \rightarrow \underbrace{[0, 255]}_{\text{range}}$$

- A color image: $f(x, y) = \begin{matrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{matrix}$

Images as discrete functions

- Images are usually **digital (discrete)**:
 - **Sample** the 2D space on a regular grid
- Represented as a matrix of integer values

pixel

j

i

62	79	23	119	120	05	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

Images as discrete functions

Cartesian coordinates

$f[n, m] =$

↑
Notation for
discrete
functions

$$\begin{bmatrix} \vdots & f[-1, 1] & f[0, 1] & f[1, 1] & \\ \dots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] & \dots \\ f[-1, -1] & f[0, -1] & f[1, -1] & & \\ & \vdots & & \ddots & \end{bmatrix}$$

What we will learn today?

- Images as functions
- **Linear systems (filters)**
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

Systems and Filters

- **Filtering:**

- Form a new image whose pixels are a combination original pixel values

Goals:

- Extract useful information from the images
 - Features (edges, corners, blobs...)
- Modify or enhance image properties:
 - super-resolution; in-painting; de-noising


De-noising



Salt and pepper noise



Super-resolution



In-painting



Bertamio et al

2D discrete-space systems (filters)

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$

$$g = \mathcal{S}[f], \quad g[n, m] = \mathcal{S}\{f[n, m]\}$$

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m]$$

Filter example #1: Moving Average

- 2D DS moving average over a 3×3 window of neighborhood

$$g[n, m] = \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l]$$

$$= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$$\frac{1}{9} \begin{matrix} & \text{h} \\ \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} \end{matrix}$$

$$(f \quad h)[m, n] = \frac{1}{9} \sum_{k, l} f[k, l] h[m - k, n - l]$$

Filter example #1: Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

Filter example #1: Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10							

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

Filter example #1: Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20						

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

Filter example #1: Moving Average

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

	0	10	20	30					

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

Filter example #1: Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30				

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

Filter example #1: Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$(f * h)[m, n] = \sum_{k, l} f[k, l] h[m - k, n - l]$$

Source: S. Seitz

Filter example #1: Moving Average

In summary:

- Replaces each pixel with an average of its neighborhood.
- Achieve smoothing effect (remove sharp features)

$$\frac{1}{9} h[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

Filter example #1: Moving Average



Filter example #2: Image Segmentation

- Image segmentation based on a simple threshold:

$$g[n, m] = \begin{cases} 255, & f[n, m] > 100 \\ 0, & \text{otherwise.} \end{cases}$$



Classification of systems

- Amplitude properties

- Linearity

- Stability

- Invertibility

- Spatial properties

- Causality

- Separability

- Memory

- Shift invariance

- Rotation invariance

Shift-invariance

If $f[n, m] \xrightarrow{\mathcal{S}} g[n, m]$ then

$$f[n - n_0, m - m_0] \xrightarrow{\mathcal{S}} g[n - n_0, m - m_0]$$

for every input image $f[n, m]$ and shifts n_0, m_0

Is the moving average system is shift invariant?

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Is the moving average system is shift invariant?

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$$f[n - n_0, m - m_0]$$

$$\xrightarrow{\mathcal{S}} \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[(n - n_0) - k, (m - m_0) - l]$$

$$= g[n - n_0, m - m_0]$$

Yes!

Linear Systems (filters)

$$f(x, y) \rightarrow \boxed{\mathcal{S}} \rightarrow g(x, y)$$

- Linear filtering:
 - Form a new image whose pixels are a weighted sum of original pixel values
 - Use the same set of weights at each point

- **S** is a linear system (function) iff it *S satisfies*

$$\mathcal{S}[\alpha f_1 + \beta f_2] = \alpha \mathcal{S}[f_1] + \beta \mathcal{S}[f_2]$$

superposition property

Linear Systems (filters)

$$f(x, y) \rightarrow \boxed{\mathcal{S}} \rightarrow g(x, y)$$

- Is the moving average a linear system?
- Is thresholding a linear system?
 - $f1[n,m] + f2[n,m] > T$
 - $f1[n,m] < T$
 - $f2[n,m] < T$ No!

LSI (linear *shift invariant*) systems

Impulse response

$$\delta_2[n, m] \rightarrow \boxed{\mathcal{S}} \rightarrow h[n, m]$$

$$\delta_2[n - k, m - l] \rightarrow \boxed{\mathcal{S} \text{ (SI)}} \rightarrow h[n - k, m - l]$$

LSI (linear *shift invariant*) systems

Example: impulse response of the 3 by 3 moving average filter:

$$h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$

$$= \begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

h

1	1	1
1	1	1
1	1	1

$\frac{1}{9}$

LSI (linear *shift invariant*) systems

An LSI system is completely specified by its impulse response.

sifting property of the delta function

$$f[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \delta_2[n - k, m - l]$$

superposition

$$\begin{array}{c} \rightarrow \boxed{\mathcal{S} \text{ LSI}} \rightarrow \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l] \\ \delta_2[n, m] \rightarrow \boxed{\mathcal{S}} \rightarrow h[n, m] \end{array}$$

Discrete convolution

$$f[n, m] * h[n, m]$$

What we will learn today?



- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



Discrete convolution (symbol: $*$)

- Fold $h[k,l]$ about origin to form $h[-k,-l]$
- Shift the folded results by n,m to form $h[n-k,m-l]$
- Multiply $h[n-k,m-l]$ by $f[k,l]$
- Sum over all k,l
- Repeat for every n,m





Discrete convolution (symbol: $*$)

- Fold $h[k,l]$ about origin to form $h[-k,-l]$
- Shift the folded results by n,m to form $h[n-k,m-l]$
- Multiply $h[n-k,m-l]$ by $f[k,l]$
- Sum over all k,l
- Repeat for every n,m



Discrete convolution (symbol: $*$)

- Fold $h[k,l]$ about origin to form $h[-k,-l]$
- Shift the folded results by n,m to form $h[n-k,m-l]$
- Multiply $h[n-k,m-l]$ by $f[k,l]$
- Sum over all k,l
- Repeat for every n,m



Convolution in 2D - examples



Original



•0	•0	•0
•0	•1	•0
•0	•0	•0



Convolution in 2D - examples



Original



•0	•0	•0
•0	•1	•0
•0	•0	•0



Filtered
(no change)

Courtesy of D Lowe

Convolution in 2D - examples




Original



•0	•0	•0
•0	•0	•1
•0	•0	•0



Convolution in 2D - examples



Original



•0	•0	•0
•0	•0	•1
•0	•0	•0



Shifted right
By 1 pixel

Convolution in 2D - examples



Original

$$\ast \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} = ?$$

Convolution in 2D - examples




Original



$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1



Blur (with a
box filter)

Courtesy of D Lowe

Convolution in 2D - examples



Original

$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix} = ?$$

(Note that filter sums to 1)

“details of the image”


$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} + \begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix}$$

A red bracket is drawn above the second 3x3 grid in the equation, spanning its width.

- What does blurring take away?



−




=



- Let's add it back:



+ a



=



Convolution in 2D – Sharpening filter



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

–

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

=



Sharpening filter: Accentuates differences with local average

Image support and edge effect

- A computer will only convolve **finite support signals**.
 - That is: images that are zero for n, m outside some rectangular region
- MATLAB's `conv2` performs 2D DS convolution of finite-support signals.





Image support and edge effect

- A computer will only convolve **finite support signals**.
- What happens at the edge?



- zero “padding”
- edge value replication
- mirror extension
- more (beyond the scope of this class)

-> Matlab conv2 uses zero-padding

What we will learn today?

- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

(Cross) correlation (symbol: $**$)

Cross correlation of two 2D signals $f[n,m]$ and $g[n,m]$

$$r_{fg}[k,l] \triangleq \sum_{n=-\infty}^{\infty} \sum_{m=-\infty}^{\infty} f[n,m] g^*[n-k, m-l]$$

$$= \sum_{n=-\infty}^{\infty} \sum_{m=-\infty}^{\infty} f[n+k, m+l] g^*[n, m], \quad k, l \in \mathbb{Z},$$

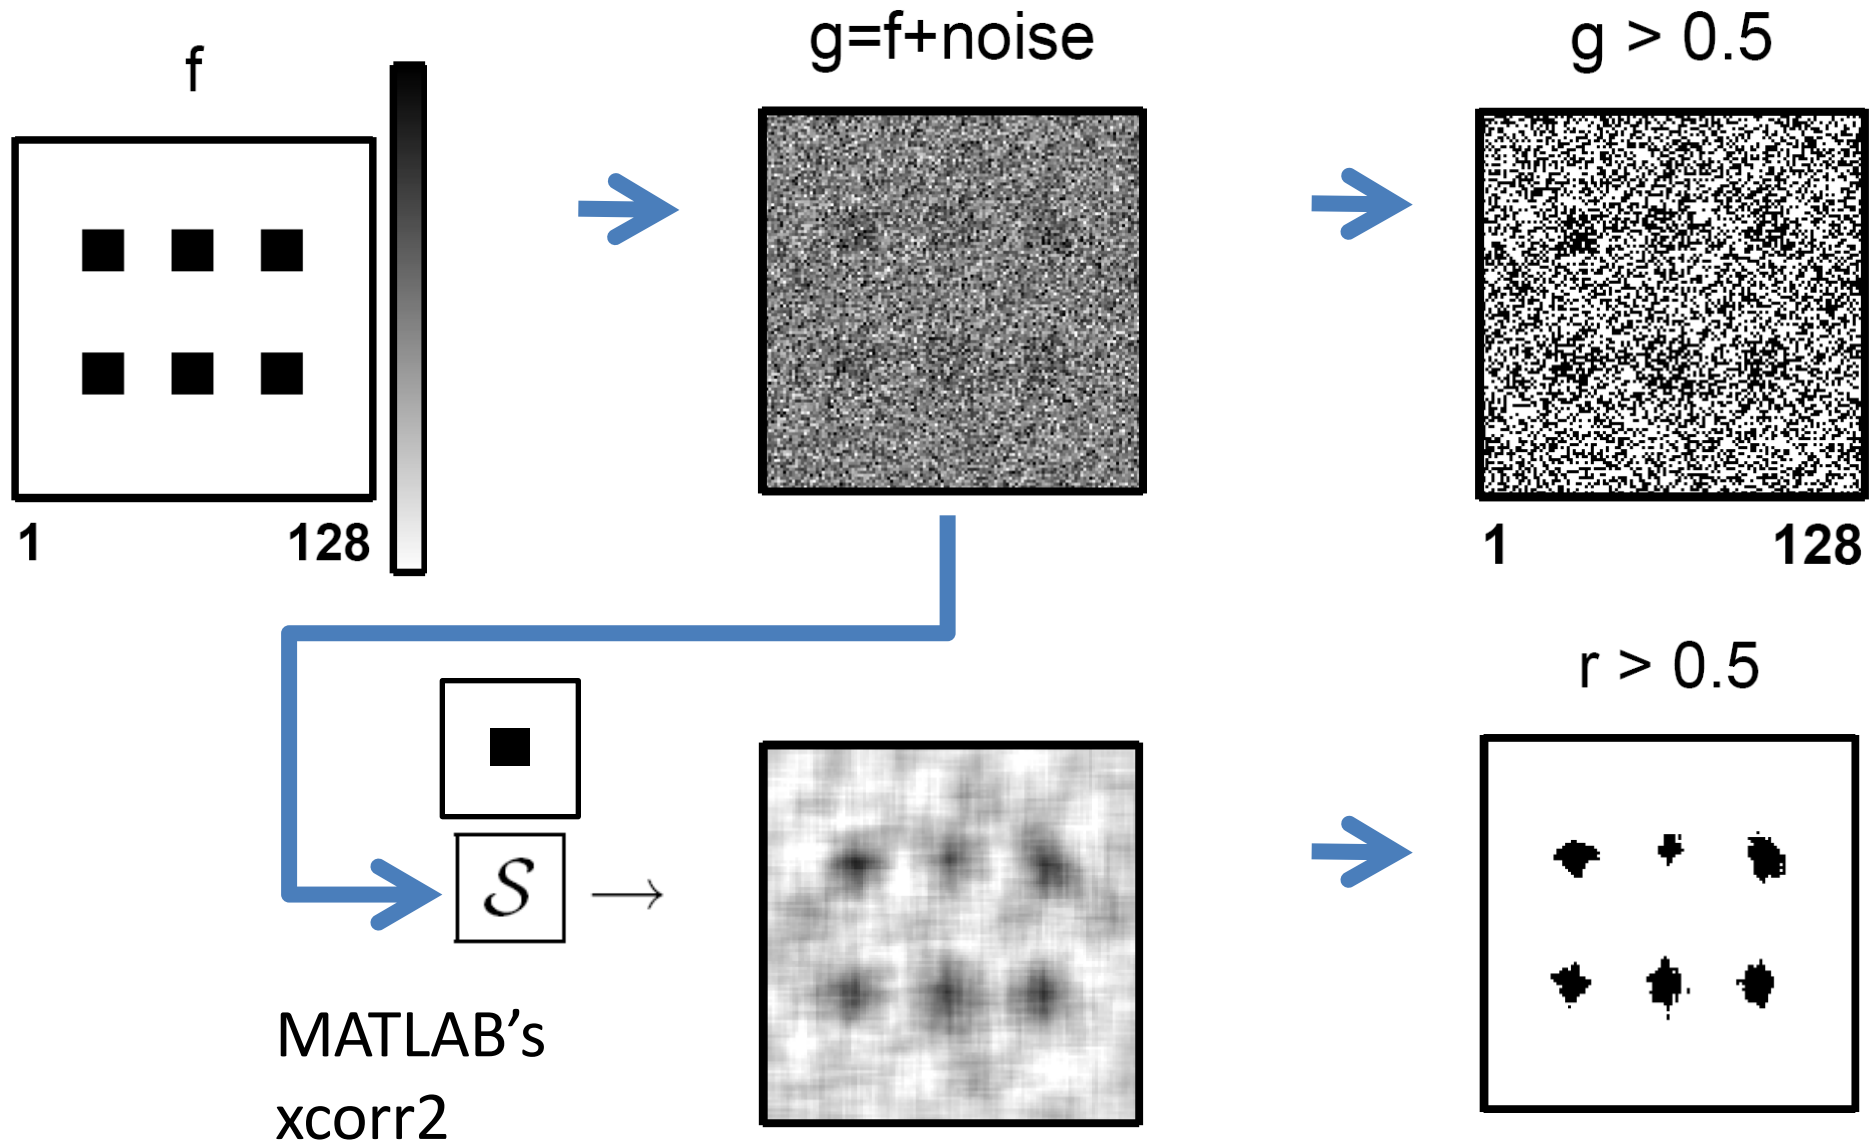
(k, l) is called the **lag**

- Equivalent to a convolution without the flip

$$r_{fg}[n, m] = f[n, m] * g^*[-n, -m]$$

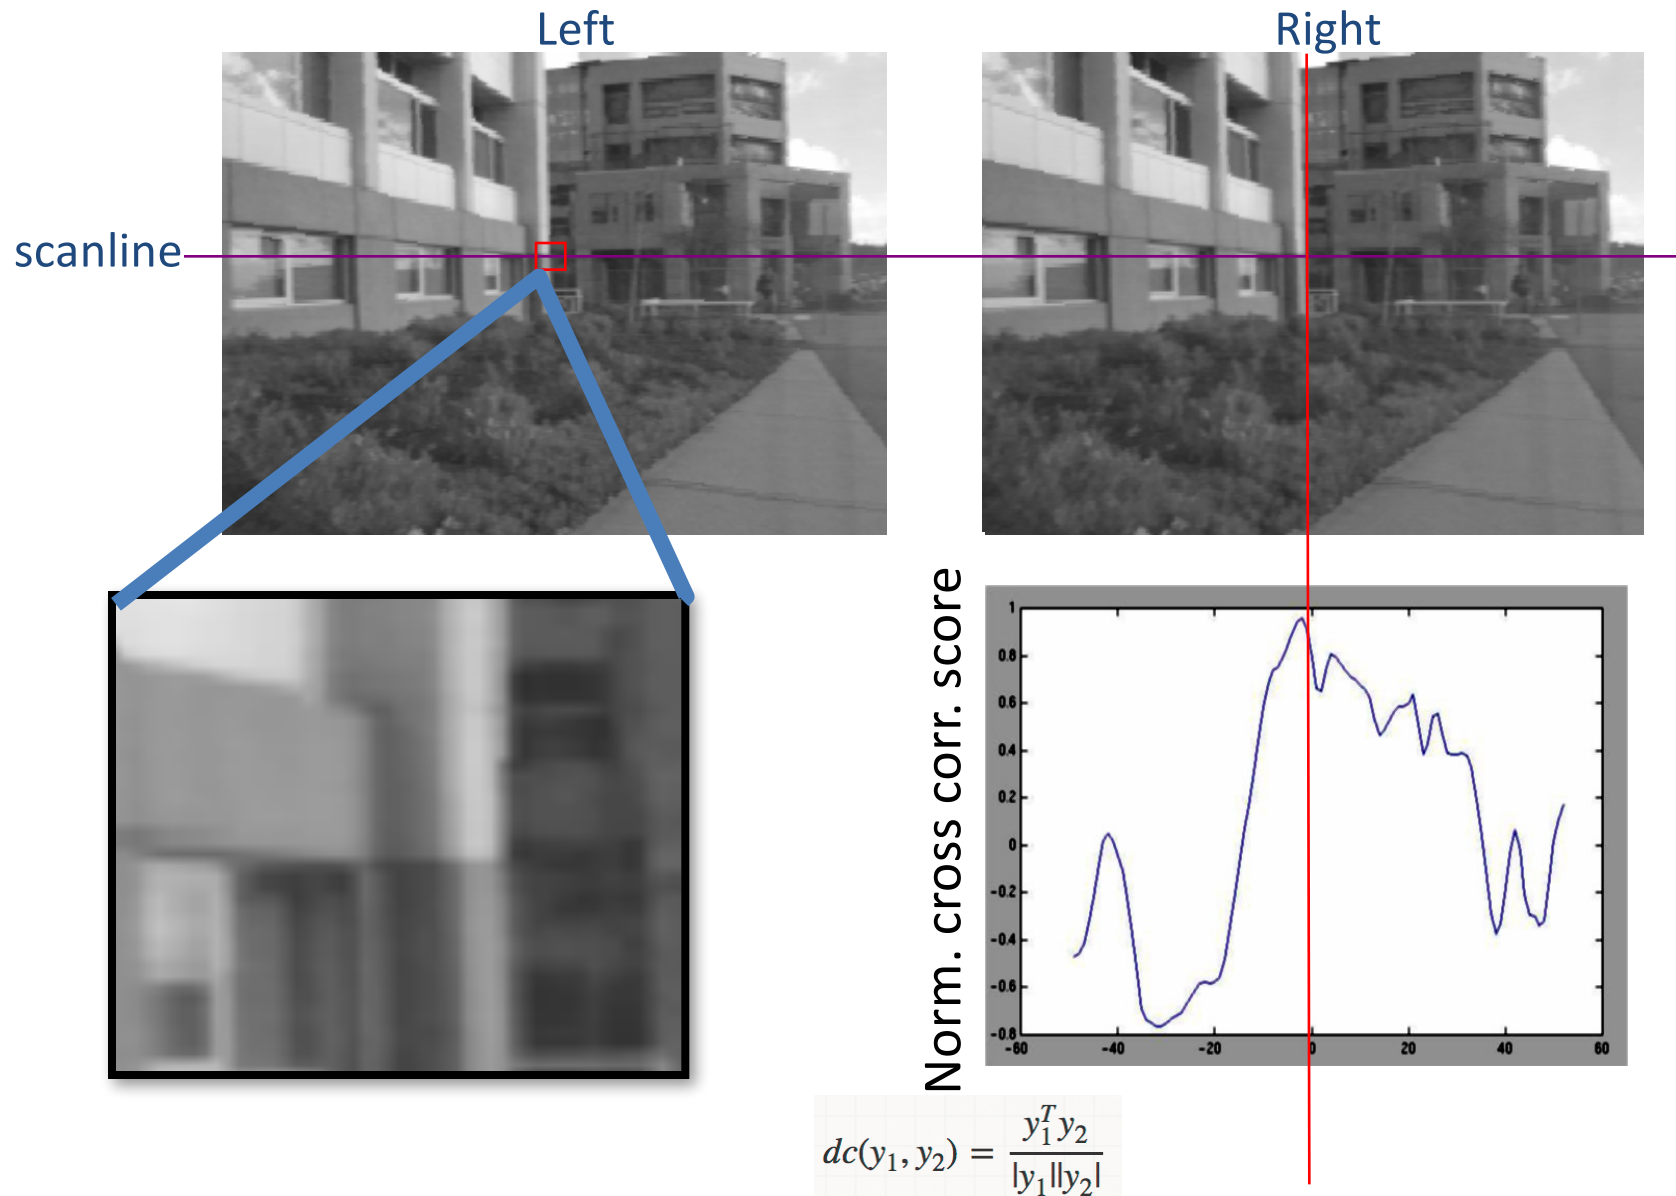
(g^* is defined as the *complex conjugate* of g . In this class, $g(n,m)$ are real numbers, hence $g^*=g$.)

(Cross) correlation – example




Courtesy of J. Fessler

(Cross) correlation – example



Convolution vs. (Cross) Correlation



Convolution vs. (Cross) Correlation

- A **convolution** is an integral that expresses the amount of overlap of one function as it is shifted over another function.
 - convolution is a filtering operation
- **Correlation** compares the *similarity of two sets of data*. Correlation computes a measure of similarity of two input signals as they are shifted by one another. The correlation result reaches a maximum at the time when the two signals match best .
 - correlation is a measure of relatedness of two signals

Cross Correlation Application: Vision system for TV remote control

- uses template matching

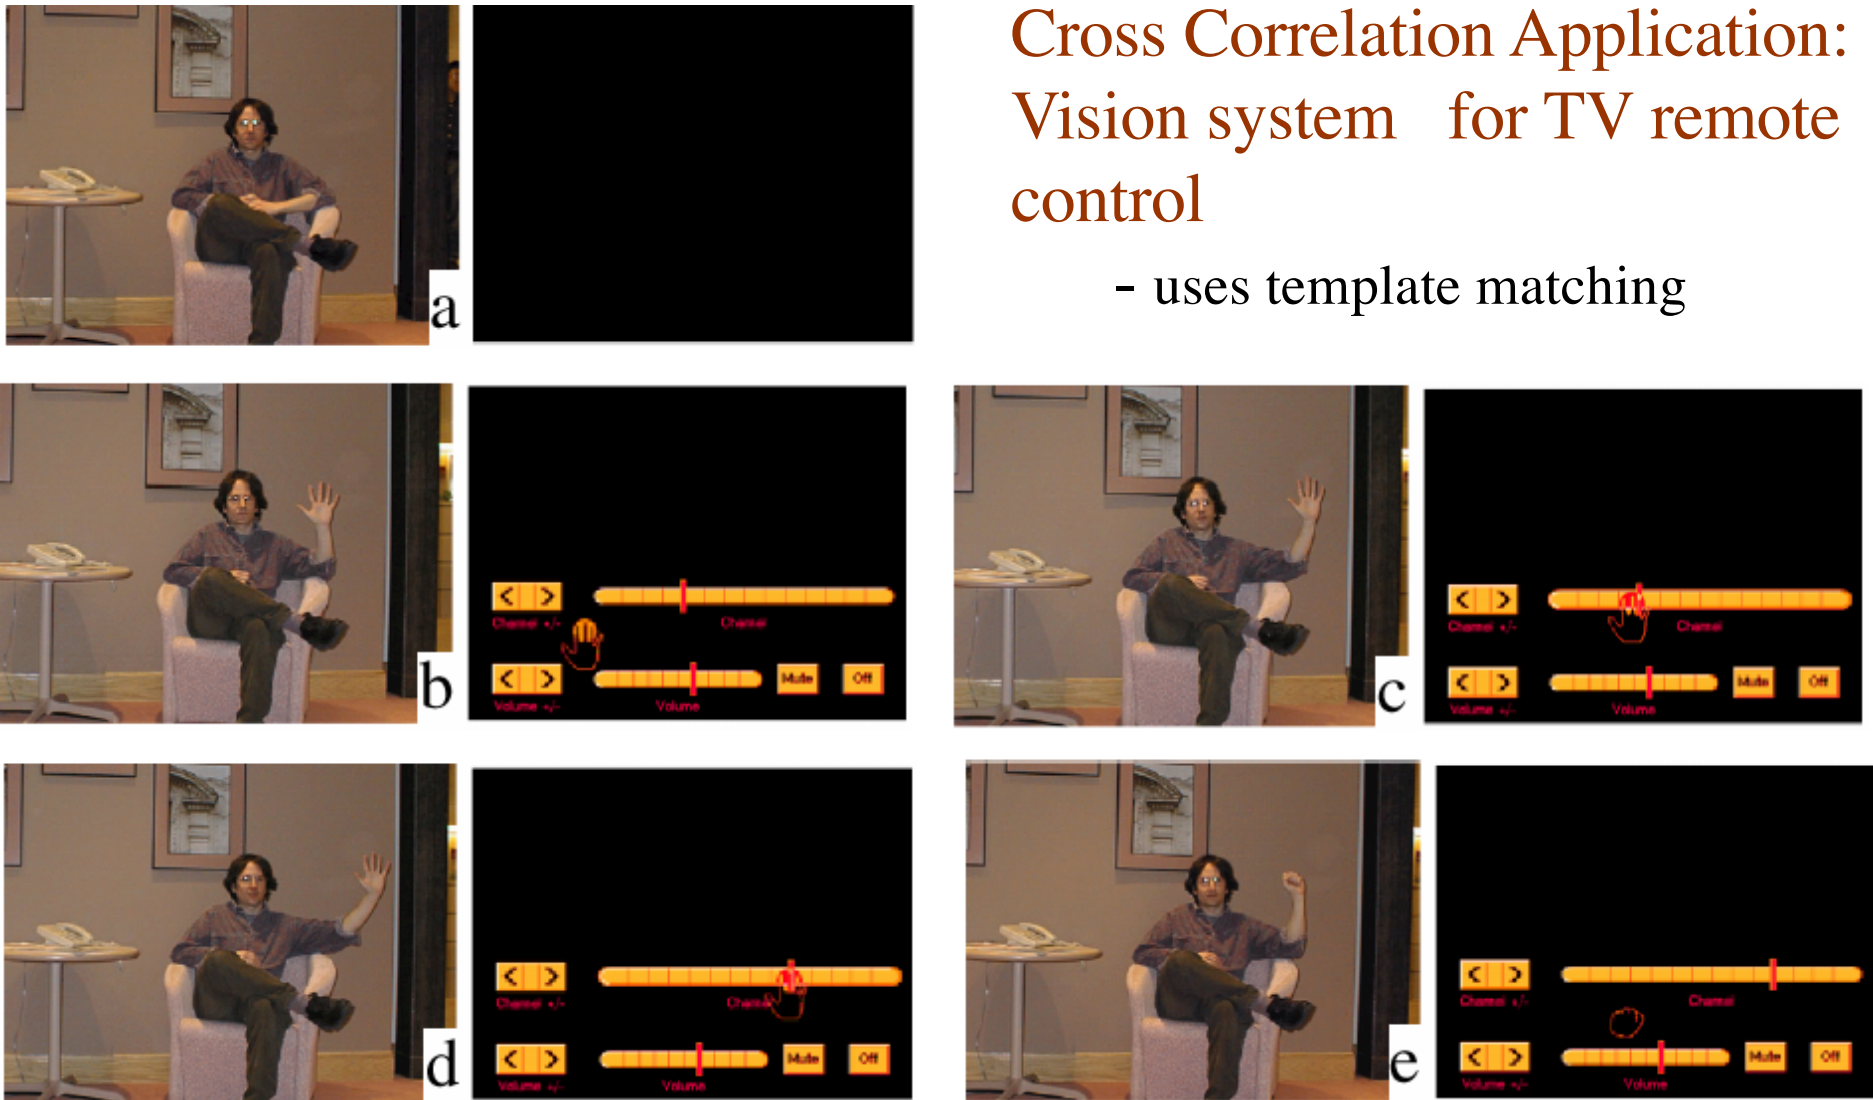


Figure from “Computer Vision for Interactive Computer Graphics,” W.Freeman et al, IEEE Computer Graphics and Applications, 1998 copyright 1998, IEEE

properties

- **Commutative property:**

$$f ** h = h ** f$$

- **Associative property:**

$$(f ** h_1) ** h_2 = f ** (h_1 ** h_2)$$

- **Distributive property:**

$$f ** (h_1 + h_2) = (f ** h_1) + (f ** h_2)$$

The order doesn't matter! $h_1 ** h_2 = h_2 ** h_1$

properties

- **Shift property:**

$$f[n, m] ** \delta_2[n - n_0, m - m_0] = f[n - n_0, m - m_0]$$

- **Shift-invariance:**

$$\begin{aligned} g[n, m] &= f[n, m] ** h[n, m] \\ \implies f[n - l_1, m - l_1] ** h[n - l_2, m - l_2] \\ &= g[n - l_1 - l_2, m - l_1 - l_2] \end{aligned}$$

What we have learned today?

- Images as functions
- Linear systems (filters)
- Convolution and correlation